

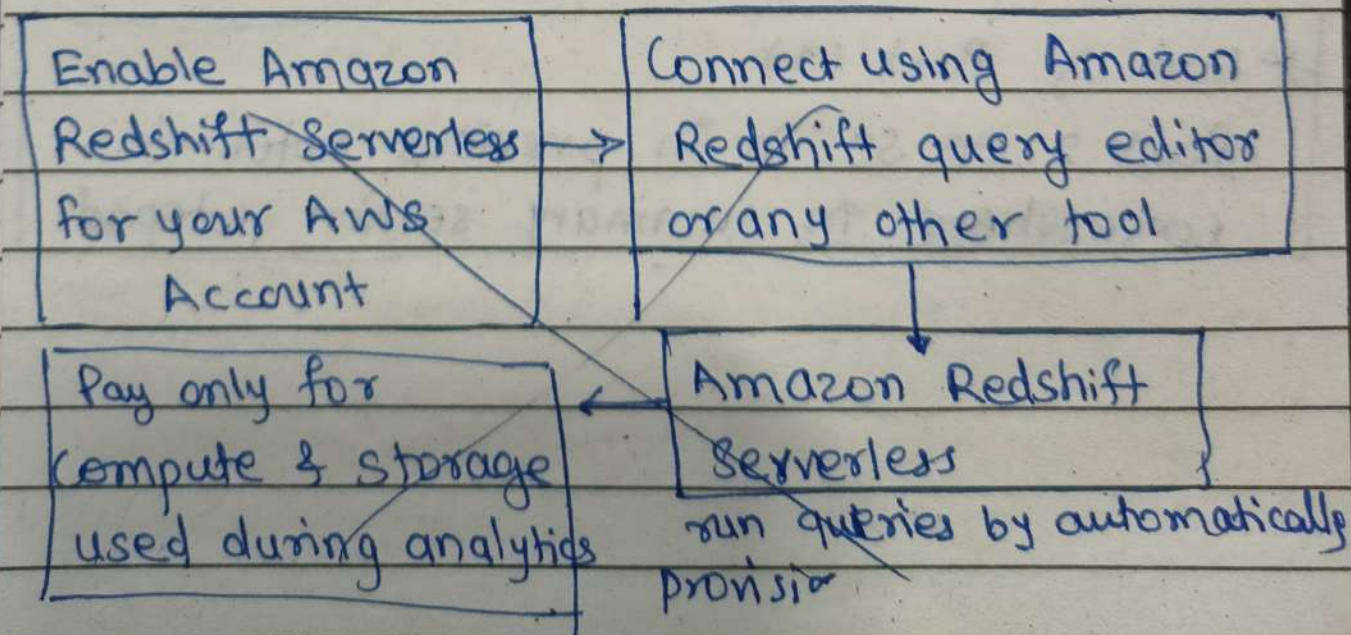
Redshift Overview

- Redshift is based on PostgreSQL, but it's not used for OLTP.
- It's OLAP - Online Analytical Processing (analytics and data warehousing)
- Load data once every hour, not every second.
- 10x better performance than other data warehouses, scale to PBs of data.
- Columnar storage of data (instead of row based)
- Massively Parallel Query Execution (MPP), highly available
- Pay as you go based on the instances provisioned.
- has SQL interface for performing the queries.

- BI tools such as AWS Quicksight or Tableau integrate with it.

Redshift Serverless

- Automatically provisions & scales data warehouse underlying capacity
- Run analytics workloads without managing data warehouse infrastructure.
- Pay only for what you use (Save costs)
- Use Cases: Reporting, dashboarding appl's, real-time analytics.



user

Connect using Amazon
Redshift query editor
or any other tool

Enable Amazon
Redshift Serverless
for your AWS Acc.

Payment

Amazon
Redshift
Serverless

Pay only for compute
& storage used
during analysis.

run queries by
automatically
provisioned & scale
capacity based on
workload

EMR Overview

- EMR stands for "Elastic MapReduce".
- EMR helps creating Hadoop clusters (Big Data) to analyze & process vast amount of data.
- The clusters can be made of hundreds of EC2 instances.
- Also supports Apache spark, HBase, Presto, Flink.....
- EMR takes care of all the provisioning & configuration
- Auto-scaling and integrated with spot instances
- Use Cases: data processing, machine learning, web indexing, big data.....

Athena Overview

- Serverless query service to perform analytics against S3 objects
- Uses standard SQL Language to query the files.
- Supports CSV, JSON, ORC, Avro, & Parquet (built on Presto)
- Pricing 5\$ per TB of data scanned.
- Use compressed or columnar data for cost-savings (less scans)
- Use-Cases: Business Intelligence / analytics / reporting, analyze & query VPC Flowlogs, ELB logs, Cloudtrail trails, etc.---
- Exam Tip :- analyze data in S3 using serverless SQL, use Athena.

Amazon Quicksight

- Serverless machine learning-powered business intelligent service to create interactive dashboards.
- Fast, automatically scalable, embeddable, with person pricing.
- Use Cases:
 - 1] Business analytics
 - 2] Business visualizations
 - 3] Perform ad-hoc analysis
 - 4] Get business insights using data.
- Integrated with RDS, Aurora, Athena, Redshift, S3.

DocumentDB (NoSQL DB)

- Aurora is an "AWS-implementation" of PostgreSQL/MySQL.....
- MongoDB is used to store, query, & index JSON data.
- Similar "deployment concepts" as Aurora.
- Fully managed, highly available with replication across 3 AZ.
- DocumentDB storage automatically grows in increments of 10 GB.
- Automatically scales to workloads with millions of requests per seconds.

Amazon Neptune

- Fully managed graph database.
- A popular graph datasets would be social network.
 - Users have friends
 - Posts have comments
 - Comments have likes from users.
 - Users share & like posts....
- Highly available across 3 AZ, with upto 15 read replicas.
- Build & run applications working with highly connected datasets - optimized for these complex & hard queries.
- Can store upto billions of relations & query the graph with milliseconds latency.
- Highly available with replications across multiple AZs.
- Great for knowledge graphs (Wikipedia), fraud detection, recommendation, social network.

Amazon Timestream

- Fully managed, fast, scalable, serverless time series database.
- Automatically scales up/down to adjust capacity.
- Store and analyse trillions of events per day
- 1000s times faster & 1/10th the cost of relational databases.
- Built-in time series analytics functions
Helps you identify patterns in your data in real-time.

Amazon QLDB

- QLDB stands for "Quantum Ledger Database".
- A Ledger is a book recording financial transactions.
- Fully managed, serverless, highly available, Replication across 3 AZ.
- Used to review history of all the changes made to your application data over time.
- Immutable system: no entry can be removed or modified, cryptographically verifiable.
- 2-3x better performance than common ledger blockchain: no decentralization component, in accordance with financial regulation rules.

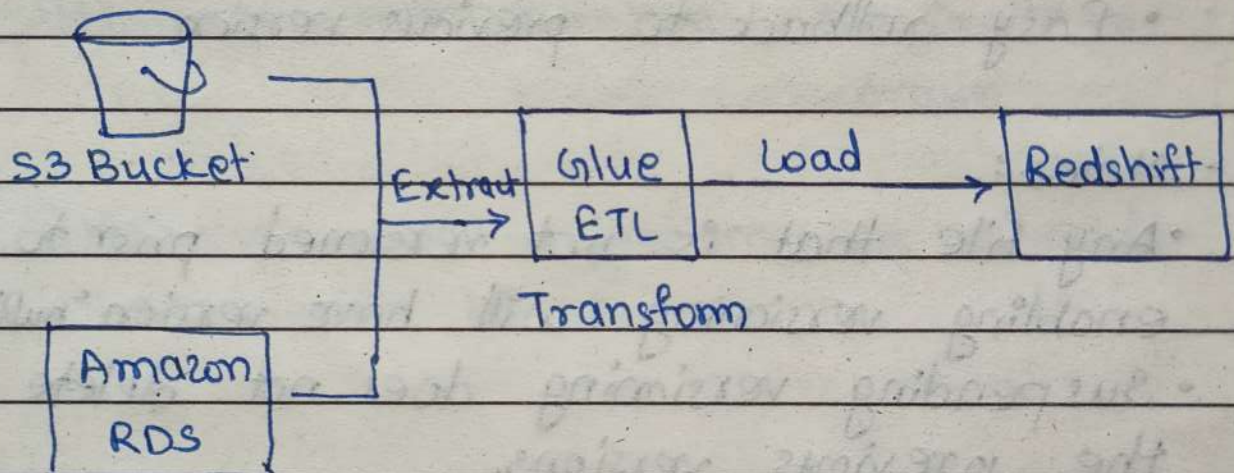
Amazon

Managed Blockchain (Decentralized)

- Blockchain makes it possible to build apps where multiple parties can execute transactions without the need of a trusted, central Authority.
- Amazon managed Blockchain is a managed service to:
 - Join public blockchain networks.
 - Or create your own scalable private network.
- Compatible with the frameworks Hyperledger Fabric & Ethereum.

AWS Glue

- Managed **extract, transform & load (ETL)** service.
- Useful to prepare & transform data for analytics.
- Fully **serverless** service.



- Glue Data Catalog: catalog of datasets
- This can be used by Athena, Redshift, EMR

DMS - Data Migration Service

- Quickly and securely migrate databases to AWS, resilient, self healing

- The source database remains available during the migration.

- Supports :

- 1] Homogeneous migrations : ex Oracle to Oracle

- 2] Heterogeneous migrations : ex Microsoft SQL Server to Aurora.

Database & Analytics Summary in AWS

- ① Relational Databases - OLTP : RDS & Aurora (SQL)
- ② Differences between Multi-AZ, Read-replicated Multi-Region
- ③ In-memory : ElastiCache
- ④ Key/Value Databases : DynamoDB (serverless) & DAX (cache for dynamoDB) (NoSQL)
(can handle millions of request per sec)
- ⑤ Warehouse - OLAP : Redshift (SQL)
(Analytics)
- ⑥
- ⑥ Hadoop cluster : EMR (Big data)
- ⑦ Athena : query data on Amazon S3
(serverless & SQL)
- ⑧ QuickSight : dashboards on your data
(serverless)

⑨ DocumentDB: "Aurora for MongoDB"
(JSON - NoSQL)

⑩ Amazon QLDB: Financial Transactions
Ledger (immutable journal, cryptographically
verifiable)

⑪ Amazon Managed Blockchain: managed
Hyperledger fabric & Ethereum blockchain
(decentralized Blockchain)

⑫ Glue: Managed ETL (extract Transform
Load) and Data Catalog Service.

⑬ Database Migration: DMS

⑭ Neptune: graph database

⑮ Timestream: time-series database